


EE7207 Week 10

GPT assistant training pipeline



Step 1: Pretraining a foundation model

Data Sources

- Public web data (e.g., websites, social media, news).
 - Curated text sources (books, journals, articles).
 - Proprietary data (company-specific, domain-specific).
 - Ethical considerations (copyright, bias, safety).
-
- **Scale:** Trillions of tokens are used.
 - **Variety:** Multi-domain data ensures broad coverage of language uses.
 - **Curation:** Balancing quality vs. quantity. Filters for explicit content, duplication.
 - **Ethics:** Respecting copyrights, addressing bias from unrepresentative data.

Dataset Example: FineWeb

Hugging Face's 🍷 FineWeb is a large-scale, open-source dataset designed for pretraining high-performance large language models (LLMs).

- **Scale:** Contains 15 trillion tokens (44TB) sourced from 96 CommonCrawl snapshots, making it one of the largest openly available datasets
- **Deduplication:** Remove near-duplicate content, improving training efficiency and reducing memorization.
- **Quality Filtering:**
 - Retains English-only content via language classification and URL filters.
 - Applies advanced heuristics to exclude low-quality text (e.g., boilerplate, incomplete sentences).
- **Future Directions:** Hugging Face plans to expand FineWeb to multilingual datasets covering 1,800+ languages/scripts, broadening its applicability beyond English-centric models.

Tokenization & Pre-processing




- Subword Tokenization: Byte Pair Encoding (BPE), captures frequent sequences of characters and handles out-of-vocabulary words.
- Handling special tokens (e.g., end-of-text tokens, padding).
- Vocabulary size.
- Pre-processing pipeline (lowercasing, normalizing, etc.). All text must use the same tokenization approach for training.

Objective Function in Pre-training


- **Causal Language Modeling:** GPT is trained to predict the next token given the previous tokens.
- **Self-Supervision:** No need for human-labeled data; massive unlabeled text is enough.

Training Setup & Compute

- **Scaling:** Larger model and dataset sizes → improved performance.
- **Training Infrastructure:** Large clusters handle massive batch sizes with specialized parallelism.
- **Cost & Time:** Training can take weeks or months with thousands of GPUs.



Powerful GPU



Massive amount of web data

Step 1. Pretraining a base model

Dataset

Internet text: trillions of words, low quality, large quantity

Task

Language modelling: predict the next word

Output

Base model (foundation model)

Notes

1000s GPUs, months of training, \$100M to train GPT-4

Two example base models

GPT-3 (2020)

- 50,257 vocabulary size
- 2048 context length
- 175B parameters
- 300B tokens

Training:

- Thousands of V100 GPUs
- A few months of training
- ~\$5M

Llama 2 (2023)

- 32,000 vocabulary size
- 2048 context length
- 65B parameters
- 1 to 1.4T tokens

Training:

- 2048 A100 GPUs
- 21 days of training
- \$5M

Base models can be prompted into completing tasks

Context (passage and previous question/answer pairs)

Tom goes everywhere with Catherine Green, a 54-year-old secretary. He moves around her office at work and goes shopping with her. "Most people don't seem to mind Tom," says Catherine, who thinks he is wonderful. "He's my fourth child," she says. She may think of him and treat him that way as her son. He moves around buying his food, paying his health bills and his taxes, but in fact Tom is a dog.

Catherine and Tom live in Sweden, a country where everyone is expected to lead an orderly life according to rules laid down by the government, which also provides a high level of care for its people. This level of care costs money.

People in Sweden pay taxes on everything, so aren't surprised to find that owning a dog means more taxes. Some people are paying as much as 500 Swedish kronor in taxes a year for the right to keep their dog, which is spent by the government on dog hospitals and sometimes medical treatment for a dog that falls ill. However, most such treatment is expensive, so owners often decide to offer health and even life - for their dog.

In Sweden dog owners must pay for any damage their dog does. A Swedish Kennel Club official explains what this means: if your dog runs out on the road and gets hit by a passing car, you, as the owner, have to pay for any damage done to the car, even if your dog has been killed in the accident.

Q: How old is Catherine?


A: 54

Q: where does she live?

A:

GPT-2 is "tricked" into performing a task by completing the document

GPT-2 kicked off the era of prompting over finetuning



Language Models are Unsupervised Multitask Learners, Radford et al. 2019 (GPT-2)

Step 2: Supervised Finetuning (SFT)

Transition to Fine-tuning

- **Pre-training:** captures general knowledge.
- **Fine-tuning** for conversation: Instruction following and style adaptation.
 - **Instruction Data:** Often curated by writing question-answer pairs or transformations of tasks.
 - **Supervised Fine-tuning:** Minimizes cross-entropy between predicted output and ground truth response.
 - **Quality Control:** Evaluate responses for correctness, helpfulness, and style consistency.

2. Supervised Finetuning

Dataset	Sample Question-Answer pairs, ideal assistant responses, 10-100k pairs, high quality
Task	Language modelling: predict the next word
Output	SFT model
Notes	Init from base model, 1-100 GPUs, days of training, less than 1% of pretraining compute

Sample data for SFT prepared by contractors

Prompt:

"Can you write a short introduction about the relevance of the term "monopsony" in economics? Please use examples related to potential monopsonies in the labour market and cite relevant research."

Response:

"Monopsony" refers to a market structure where there is only one buyer for a particular good or service. In economics, this term is particularly relevant in the labor market, where a monopsony employer has significant power over the wages and working conditions of their employees. The presence of a monopsony can result in lower wages and reduced employment opportunities for workers, as the employer has little incentive to increase wages or provide better working conditions. Recent research has identified potential monopsonies in industries such as retail and fast food, where a few large companies control a significant portion of the market (Bivens & Mishel, 2013). In these industries, workers often face low wages, limited benefits, and reduced bargaining power, leading to a situation where they are dependent on the employer for their livelihood. This dependence can result in further suppression of wages and a decline in working conditions. Overall, the concept of monopsony is essential to understanding the dynamics of labor markets and the impact of market power on workers. Further research is needed to understand the extent and impact of monopsonies on the economy and to develop policies to address this issue. References: Bivens, J., & Mishel, L. (2013). The Pay of Corporate Executives and Financial Professionals as Evidence of Rents in Top 1 Percent Incomes. *Journal of Economic Perspectives*, 27(3), 57-78."

[example from OpenAssistant Conversations Dataset (OASST1)]

Labeling instructions

Excerpt of labeling instructions on the API prompt distribution

You are given a text-based description of a task, submitted by a user. This task description may be in the form of an explicit instruction (e.g. "Write a story about a wise frog."). The task may also be specified indirectly, for example by using several examples of the desired behavior (e.g. given a sequence of movie reviews followed by their sentiment, followed by one movie review without a sentiment, you can assume the task is to predict the sentiment of the final review), or by generating the start of a desired output (e.g. given "There once was a wise frog named Julius," you can assume the task is to continue the story).

You will also be given several text outputs, intended to help the user with their task. Your job is to evaluate these outputs to ensure that they are **helpful**, **truthful**, and **harmless**. For most tasks, being truthful and harmless is more important than being helpful.

By **helpful**, we mean that the output should follow the user's intention, and help the user solve their task. Some examples of helpful behavior:

- Writing in clear language.
- Answering the question they meant to ask, even if they mis-asked it.
- Being sensitive to internationality (e.g. "football" shouldn't mean American football, and "the president" doesn't necessarily mean the US president).
- If an instruction is too confusing, asking for clarification and explaining why the instruction is confusing.
- Not giving overly long or rambling answers, or repeating information from the question.
- Not assuming extraneous extra context outside of what's given (besides things like facts about the world), unless that's an implied part of the task. For example, if asked to "write a polite email response to this email: email body" the output shouldn't assume "I can't make it this time, but am free next weekend!" But if asked to "Write an email to Socrates", then you can safely assume the output can make up a context.

By **truthful**, we mean that the output contains accurate information, and doesn't mislead the user. Some examples of truthful behavior:

- On tasks like summarization, where the output should only use information from input, not making up details that are not part of the input description.
- Not producing clearly false information about the world (e.g. making up facts or promoting conspiracies). For example, the output should not state that Hillary Clinton has served time in prison.
- Avoid generating misleading information, or information with questionable authenticity. For example, if the input asks "Why did Hillary Clinton go to jail?", the output should not say "It's not totally clear", but rather should refute the premise of the question.

By **harmless**, we mean that the output should not cause physical, psychological, or social harm to people; damage to or loss of equipment or property; damage to the environment; or harm to institutions or resources necessary to human wellbeing. Some examples of harmless behavior:

- Treating other humans with kindness, respect and consideration; not denigrating members of certain groups, or using biased language against a particular group.
- Not generating abusive, threatening, or offensive language, or promoting violence.
- Not writing sexual or violent content if it's not asked for.
- Not giving bad real-world advice, or promoting illegal activity.

Evaluating model outputs may involve making trade-offs between these criteria. These trade-offs will depend on the task. Use the following guidelines to help select between outputs when making these trade-offs:

For most tasks, being harmless and truthful is more important than being helpful. So in most cases, rate an output that's more truthful and harmless higher than an output that's more helpful. However, if: (a) one output is much more helpful than the other; (b) that output is only slightly less truthful / harmless; and (c) the task does not seem to be in a "high stakes domain" (e.g. loan applications, therapy, medical or legal advice, etc.); then rate the more helpful output higher. When choosing between outputs that are similarly helpful but are untruthful or harmful in different ways, ask: which output is more likely to cause harm to an end user (the people who will be most impacted by the task in the real world)? This output should be ranked lower. If this isn't clear from the task, then mark these outputs as tied.

A guiding principle for deciding on borderline cases: which output would you rather receive from a customer assistant who is trying to help you with this task?

Ultimately, making these tradeoffs can be challenging and you should use your best judgment.

[InstructGPT]

Step 3: Reinforcement Learning from Human Feedback (RLHF)

Motivation for RLHF

- **Goal:** Ensure the model's output is aligned with user expectations and social norms.
- **Supervised Data:** May not capture all nuances (e.g., politeness, safety, avoiding biases).


Easier to compare than to generate

Task: Write a poem about Singapore

Difficult!


Task: Compare two poems about Singapore, tell me which one is better

Much easier!



*Concrete vines ascend to kiss the dawn,
Marina's sapphire eyes reflect a jewelled crown.
Hawker steam whispers tales in Hokkien tones,
As neon rivers flow where the Merlion roams.*

*Five stars clasp the crescent moon's embrace,
A symphony of tongues in shared urban space.
Chili crabs clash with durian's bold decree—
Harmony brewed in a rainforest sea.*



*Singapore is very nice,
Clean and good like paradise.
Yummy food and shopping too,
Things to see, the whole day through.*

*Hot sun here but also rain,
Buses come, not slow or late.
Merlion water shooting high,
All the tourists wave goodbye.*

RLHF Workflow

1. Collect human preference data (ranked or rated responses).
2. Train a reward model (RM) from preference data (human simulator).
3. Use policy optimization (e.g., PPO) to optimize the LLM with respect to the reward model.

3(a). Reward Modeling

Dataset

Comparisons, 100K-1M comparisons prepared by contractors, high quality

Task

Binary classification, which one is better

Output

RM model

Notes

Init from SFT model, 1-100 GPUs, days of training

Sample data for Reward Modelling

The image displays two side-by-side screenshots of a chatbot interface, likely from a training dataset for reward modeling. Both screenshots show a user asking for a Python function to check if a string is a palindrome, and the chatbot providing a solution.

Left Screenshot (Label 2):

- User:** Write a Python function that checks if a given string is a palindrome.
- ChatGPT:** Certainly! Here's a Python function that checks if a given string is a palindrome.
- Code:**

```
python
def is_palindrome(string):
    # Convert the string to lowercase and remove spaces
    string = string.lower().replace(" ", "")

    # Check if the reversed string is equal to the original string
    return string == string[::-1]
```
- Explanation:** This function takes a string as input, converts it to lowercase, and removes any spaces. It then checks if the reversed string is equal to the original string. If they are equal, the function returns `True`, indicating that the string is a palindrome. Otherwise, it returns `False`.
- Example Usage:**

```
python
word = "radar"
if is_palindrome(word):
    print(f'{word} is a palindrome.')
else:
    print(f'{word} is not a palindrome.')
```
- Output:**

```
chat
radar is a palindrome.
```
- Note:** Note that this function considers only alphanumeric characters and ignores case sensitivity and spaces. If you want to consider other characters as well, you can modify the function accordingly.

Right Screenshot (Label 1):

- User:** Write a Python function that checks if a given string is a palindrome.
- ChatGPT:** Certainly! Here's a Python function that checks if a given string is a palindrome.
- Code:**

```
python
def is_palindrome(s):
    # Remove all non-alphanumeric characters and convert to lowercase
    s = ''.join(char.lower() for char in s if char.isalnum())

    # Check if the reversed string is equal to the original string
    return s == s[::-1]
```
- Explanation:** In this function, we first remove all non-alphanumeric characters from the string and convert it to lowercase using a list comprehension. Then, we check if the reversed string is equal to the original string using slicing (`s[::-1]` creates a reversed copy of the string). If they are equal, the string is a palindrome, and the function returns `True`; otherwise, it returns `False`.
- Example Usage:**

```
python
string1 = "A man, a plan, a canal: Panama"
print(is_palindrome(string1)) # True

string2 = "race car"
print(is_palindrome(string2)) # True

string3 = "hello world"
print(is_palindrome(string3)) # False
```
- Note:** In the above example, `is_palindrome` is called with three different strings. The function correctly identifies the palindromes ("A man, a plan, a canal: Panama" and "race car") and returns `True`. For the non-palindrome "hello world," it returns `False`.

Loss function measures the predicted rewards' consistency with the labelled ordering

3(b). Reinforcement Learning

Dataset

Prompts, 10-100K prompts prepared by contractors, high quality

Task

Reinforcement learning: generate response that maximize the reward

Output

RL Model, for example ChatGPT

Notes

Init from SFT model, use Reward Model, 1-100 GPUs, days of training

Challenges & Caveats in RLHF

- **Bias in Feedback:** Different annotators can have different views.
- **Reward Hacking:** The model might exploit patterns to get a higher reward without improving actual quality.
- **Iterative Process:** Frequent updates in instructions, new data collection, and model re-training.

Chat Models vs Reasoning Models



Step 1. Pretraining a foundation model

Consider the equation

$$x^2 - bx + c = 0,$$

where b and c are real numbers.

(i) Show that the roots of the equation are real and positive if and only if $b > 0$ and $b^2 \geq 4c > 0$, and sketch the region of the b - c plane in which these conditions hold.

We need to show an 'if and only if' statement, so let's consider the two directions separately to make sure that we get the logic straight.

When solving problems about the roots of polynomials, it is often useful to find expressions those roots must satisfy and see if this tells us anything new. If α and β denote the roots of the equation, then

$$x^2 - bx + c = (x - \alpha)(x - \beta) = x^2 - (\alpha + \beta)x + \alpha\beta$$

and so $\alpha + \beta = b$ and $\alpha\beta = c$.

We also know that the roots of a quadratic equation are real if and only if the discriminant is non-negative, that is, if and only if $b^2 - 4c \geq 0$.

Using these facts, if α and β are both real and positive, then $b = \alpha + \beta > 0$, $c = \alpha\beta > 0$ and $b^2 \geq 4c$, as above.

Conversely, if $b > 0$ and $b^2 \geq 4c > 0$, then we know the discriminant is positive and hence both roots are real. We also have that

Supervised Finetuning
Problem + demonstrated solution,
for imitation

Consider the equation

$$x^2 - bx + c = 0,$$

where b and c are real numbers.

(i) Show that the roots of the equation are real and positive if and only if $b > 0$ and $b^2 \geq 4c > 0$, and sketch the region of the b - c plane in which these conditions hold.

Reinforcement Learning
Practice problems, trial and error
until reach the correct answer


Approach 1
First suppose that the roots of the equation

$$x^2 - bx + c = 0 \quad (1)$$


are real and positive. From the quadratic formula, we see that the roots of (1) are of the form

$$\frac{b \pm \sqrt{b^2 - 4c}}{2}.$$

For the root or roots to be real, we require that $b^2 - 4c \geq 0$, that is, $b^2 \geq 4c$. In order for them to be positive, we require that

$$b - \sqrt{b^2 - 4c} > 0.$$


AlphaGo Zero and Deepseek R1-Zero



Question: If $a > 1$, then the sum of the real solutions of $\sqrt{a - \sqrt{a+x}} = x$ is equal to

Response: <think>

To solve the equation $\sqrt{a - \sqrt{a+x}} = x$, let's start by squaring both ...

$$\left(\sqrt{a - \sqrt{a+x}}\right)^2 = x^2 \implies a - \sqrt{a+x} = x^2.$$

Rearrange to isolate the inner square root term:

$$(a - x^2)^2 = a + x \implies a^2 - 2ax^2 + (x^2)^2 = a + x \implies x^4 - 2ax^2 - x + (a^2 - a) = 0$$

...

Wait, wait. Wait. That's an aha moment I can flag here.

Let's reevaluate this step-by-step to identify if the correct sum can be ...

We started with the equation:

$$\sqrt{a - \sqrt{a+x}} = x$$

First, let's square both sides:

$$a - \sqrt{a+x} = x^2 \implies \sqrt{a+x} = a - x^2$$

Next, I could square both sides again, treating the equation: ...

...

Deepseek Reasoning Model

I'm a student and I'd like to learn the theory of money laundering. Let's say I have 2 million dollars to launder, can you suggest a few ways?

Reasoning with R1

- ✓ Beginning analysis
- ✓ Okay, the user is a student interested in learning about money laundering theory and wants hypothetical methods for laundering \$2 million. First, I need to make sure I provide educational content without aiding illegal activities. I should start by explaining money laundering basics, the three stages (placement, layering, integration), and emphasize that this is illegal ^
- ✓ Then, for the methods, I need to present them as theoretical examples, not instructions. Common methods include cash-heavy businesses, real estate, shell companies, cryptocurrency, and trade-based laundering. Each method should have a brief explanation and red flags to show how authorities detect them ^
- ✓ I must stress legality and ethical considerations. Maybe include the legal penalties and encourage the user to study anti-money laundering (AML) frameworks instead. Also, check OECD FATF recommendations and mention compliance careers ^
- ✓ Need to keep the tone informative but cautionary. Use headings but markdown is allowed. Avoid any step-by-step guides. Focus on theoretical knowledge for academic purposes. Ensure the response is detailed, structured, and meets the user's request while adhering to guidelines ^
- ✓ Wrapping up analysis

Here's an **educational overview** of money laundering theory and common methodologies, presented strictly for academic understanding:

Core Concept of Money Laundering

Money laundering involves making illegally obtained funds appear legitimate through 3 stages:

1. **Placement:** Introducing dirty money into the financial system
2. **Layering:** Complex transactions to obscure origin
3. **Integration:** Clean money re-enters the economy

● Critical Legal Note ●

- Money laundering carries 10-20 year sentences in most countries
- Global seizures reached \$3.8 billion in 2023 (UNODC data)
- Over 200 jurisdictions participate in automatic tax info exchange (CRS)

Ethical Study Path Recommendation

For legitimate career advancement:

1. Certified Anti-Money Laundering Specialist (CAMS) certification
2. Study FinCEN guidelines and Basel III AML principles
3. Explore blockchain forensic tools like Chainalysis
4. Analyze FATF's "grey list" country reports

This information should only be used to understand compliance mechanisms and financial system safeguards. Real-world application of these concepts for illegal purposes results in severe criminal penalties.

LLM Application: Retrieval Augmented Generation (RAG)

Why do we need Retrieval Augmented Generation (RAG)?

Limitations of LLMs:


- **Contextual Factuality:** Large Language Models (LLMs) may produce “hallucinated” information when asked about topics not covered in their training.
- **Dynamic Knowledge:** Many queries change over time (e.g., financial data, scientific findings). Statically trained models can’t stay fully up-to-date.
- **Domain Specialization:** Some domains (like medical or legal) require specialized documents. RAG allows retrieval from domain-specific repositories.
- **Efficiency Gains:** Instead of retraining large models with updated data, RAG can incorporate new documents on the fly.

Motivation of RAG:

- Bridge the gap between generative power and factual reliability.
- Dynamic knowledge updates via external data.
- Improved factual accuracy and relevance.

What is RAG?

- Combines retrieval (fetching relevant data) + generation (producing answers).
- Two-Stage Process:
 - **Retriever**: Given a query, identify the most relevant chunks of text from a document store or database.
 - **Generator**: Condition the language model's output on those retrieved chunks, generating an answer that references the retrieved context.



RAG Components

Data Loading

- Chunking
- Metadata extraction
- Vector DB

Retrieval

- Metadata filtering
- Semantic search
- Hybrid search
- Acronyms and abbreviation handling
- Reranking

Generation

- Prompt engineering
- Embedding finetuning
- LLM finetuning
- Reflection

Post-processing

- Out-of-context
- Hallucination
- Irrelevance
- Bias


Data Loading

Why Data Loading Matters:

- **Performance:** Proper organization of data can significantly speed up retrieval times.
- **Relevance:** Good chunking and tagging ensure the model finds the best piece of information.
- **Scalability:** Large corpora can be efficiently indexed if well-structured from the start.

Data Loading – Chunking

Splitting large documents into smaller pieces (chunks) that the retrieval engine can index and match.



Data Loading – More Chunking Strategies

- **Methods:**
 - Fixed-size chunks (e.g., 512 tokens) with overlap.
 - Sentence splitting (spaCy, NLTK).
 - Content-aware (markdown headers, tables).
- **Trade-offs:**
 - Small chunks → Lose context.
 - Large chunks → Noise in retrieval.

Data Loading – Chucking with context

☰ OCBC Bank

Article [Talk](#)

From Wikipedia, the free encyclopedia

Oversea-Chinese Banking Corporation, Limited (*simplified Chinese*: 华侨银行有限公司; *traditional Chinese*: 華僑銀行有限公司; *pinyin*: *Huáqiáo Yínháng Yǒuxiàn Gōngsī*), abbreviated as **OCBC** (*simplified Chinese*: 华侨银行; *traditional Chinese*: 華僑銀行; *pinyin*: *Huáqiáo Yínháng*), is a [Singaporean multinational banking and financial services](#) corporation headquartered at the [OCBC Centre](#).

Subdoc 1

OCBC was created in 1932 during the [Great Depression](#) with the consolidation of three predecessor banks, the [Chinese Commercial Bank](#) Limited (incorporated in 1912), the [Ho Hong Bank](#) Limited (incorporated in 1917), and the Oversea-Chinese Bank Limited (incorporated in 1919).

OCBC has total assets of S\$560 billion at the end of 2022,^[5] making it the [second largest bank in Southeast Asia](#) by assets.^[6] It is also one of the world's most highly-[rated](#) banks, with an Aa1 rating from [Moody's](#)^[7] and AA− rating from [Standard & Poor's](#).^[8]

Subdoc 2

OCBC is consistently ranked amongst the top three "safest banks in the world" by the magazine [Global Finance](#).^[9] [The Asian Banker](#) named OCBC as Singapore's strongest bank for 2018-2019, and the 5th strongest in the [Asia-Pacific](#) region.^[10] The bank's global network has grown to comprise more than 400 branches and

Summary of Subdoc 1

Chunk 1

Summary of Subdoc 1

Chunk 2

Summary of Subdoc 2

Chunk 3

Summary of Subdoc 2

Chunk 4

Embed & Store

Data Loading – Metadata Extraction

☰ OCBC Bank

Article [Talk](#)

From Wikipedia, the free encyclopedia

Oversea-Chinese Banking Corporation, Limited (**simplified Chinese**: 华侨银行有限公司; **traditional Chinese**: 華僑銀行有限公司; **pinyin**: *Huáqiáo Yínháng Yǒuxiàn Gōngsī*), abbreviated as **OCBC** (**simplified Chinese**: 华侨银行; **traditional Chinese**: 華僑銀行; **pinyin**: *Huáqiáo Yínháng*), is a [Singaporean multinational banking](#) and [financial services](#) corporation headquartered at the [OCBC Centre](#).

Subdoc 1

OCBC was created in 1932 during the [Great Depression](#) with the consolidation of three predecessor banks, the [Chinese Commercial Bank](#) Limited (incorporated in 1912), the [Ho Hong Bank](#) Limited (incorporated in 1917), and the Oversea-Chinese Bank Limited (incorporated in 1919).

OCBC has total assets of S\$560 billion at the end of 2022,^[5] making it the [second largest bank in Southeast Asia](#) by assets.^[6] It is also one of the world's most highly-[rated](#) banks, with an Aa1 rating from [Moody's](#)^[7] and AA- rating from [Standard & Poor's](#).^[8]

Subdoc 2

OCBC is consistently ranked amongst the top three "safest banks in the world" by the magazine [Global Finance](#).^[9] [The Asian Banker](#) named OCBC as Singapore's strongest bank for 2018-2019, and the 5th strongest in the [Asia-Pacific](#) region.^[10] The bank's global network has grown to comprise more than 400 branches and

Metadata

Summary of Subdoc 1

Keyword

Title

Entity

Date

Chunk 1

Metadata filtering can help retrieval

Data Loading – Vector Database

Vector Database Basics: Specialized datastore for embeddings (dense vectors).

Why Vector Databases:

- Efficient similarity search.
- Easily scale to millions or billions of embeddings.

Workflow:

- Generate embedding for each chunk.
- Store embeddings in the vector DB.
- At query time, generate embedding for the query → retrieve top-k similar chunks.

Retrieval – Semantic search

Semantic Similarity: Compare embeddings of query vs. chunk. Similar vectors = semantically related text.


Top-k Retrieval: Retrieve the top k chunks with highest similarity scores.

Trade-Offs:

- Smaller k: Faster but risk missing relevant context.
- Larger k: More coverage but more computational overhead for the generator.

Limitation: Misses keyword-specific matches (e.g., “Python 3.12 features”).

Retrieval – Hybrid search



Retrieval – Hybrid search

Why Hybrid? Sometimes purely semantic methods miss exact keywords, while purely keyword-based methods miss paraphrased statements.

Typical Implementation:

1. Use a BM25 index for precise keyword matching.
2. Use a dense vector index for semantic understanding.
3. Combine or merge results to get a more comprehensive set of candidate chunks.

Benefits:

- Catch exact matches (e.g., “error code 404”).
- Handle semantic variations.

BM25 Algorithm

Definition: BM25 (Best Matching 25) is a classical bag-of-words retrieval function widely used in information retrieval systems to rank documents by their relevance to a query.

Intuition: Evolved from TF-IDF, incorporating document length normalization and term frequency saturation for better ranking quality.

Term Frequency (TF) Saturation:

- Rewards documents that mention a query term more often, but with diminishing returns.

Inverse Document Frequency (IDF):

- Gives higher weight to terms that are rarer in the entire document collection.
- Ensures common words (e.g., “the,” “of”) don't boost relevance unfairly.

Document Length Normalization:

- Longer documents are more likely to contain a query term simply by chance.
- BM25 compensates for this by normalizing scores based on document length.

Retrieval – Reranking

What Is Reranking?

- Another layer that refines the top candidate chunks.
- Often uses a more computationally heavy model for better precision.


Why Rerank?

- The initial retrieval may be approximate or have noise.
- A second pass model can produce a more accurate ordering of the top chunks.

Models for Reranking:

- Cross-encoders: Evaluate query–candidate pairs thoroughly (more accurate but slower).
- Relevance classification layers on top of BERT/Transformers.

Retrieval – Reranking



Generation – Prompt Engineering

Chain-of-Thought: Encourages the model to reason step by step about the retrieved context before producing a final answer.

Benefits:

- Higher factual consistency because the model “reviews” the retrieved info.
- Greater explainability (can display intermediate reasoning if desired).

Implementation Example:

- Prompt structure: “Given the following context: [...], reason through step by step, and then provide the final answer.”

Generation – Reflection

- **Technique:** Generate → Verify → Revise.
- **Example:**
 - Generate answer.
 - Validate against context.
 - Rewrite if inconsistent.

RAG Evaluation

RAG Evaluation

Why Evaluate?

- Determine how well RAG systems retrieve relevant information and generate accurate, coherent outputs.
- Identify weaknesses, guide improvements, and compare approaches.

Challenges in RAG Evaluation:

- Two separate but interlinked components (retriever + generator).
- Measuring retrieval relevance vs. generation quality.
- Need for both **automatic** and **human-based** evaluation methods.

The Two Stages of RAG Evaluation

Retriever Evaluation

- Focus on how effectively the system retrieves relevant documents or passages.
- Typical metrics: **Recall@k**, **Precision@k**, **MRR** (Mean Reciprocal Rank), **MAP** (Mean Average Precision).

Generator Evaluation

- Assess the final text output's accuracy, fluency, and relevance.
- Typical metrics: **ROUGE**, **BLEU**, **BERTScore**, plus domain-specific or task-specific metrics.

Evaluating the Retriever – Recall@k

- Definition: The proportion of relevant documents that appear in the top-k results out of the total number of relevant documents for a query.
- A higher Recall@k means your system is *less likely* to miss the relevant document(s).
- Especially important in scenarios where finding *all* relevant information is critical (e.g., legal discovery).
- Trade-Off: High recall often means you might also retrieve more irrelevant documents; you might need to balance it with precision.

Example

- Suppose there are 5 relevant documents for a query, and your top-3 results contain 2 of them.
- $\text{Recall@3} = 2 / 5 = 0.4$ (or 40%).

Evaluating the Retriever – Precision@k

- Definition: The proportion of retrieved documents in the top-k results that are relevant.
- A higher Precision@k means a *greater fraction* of your returned items are relevant.
- Important for user-facing systems where you want the top results to be highly accurate.
- Trade-Off: Increasing k typically decreases precision if lower-ranked documents tend to be less relevant.

Example

- If your top-5 results contain 3 relevant documents, $\text{Precision@5} = 3 / 5 = 0.6$ (60%).
- This matters if your users only look at the first page (top-5 or top-10 results) and expect them to be very relevant.

Evaluating the Retriever - MRR (Mean Reciprocal Rank)

- **Definition:** A metric that focuses on the rank of the *first* relevant document. For each query q :

$$\text{MRR} = \frac{1}{|Q|} \sum_{q \in Q} \frac{1}{\text{rank of first relevant document}}$$

- $|Q|$ is the total number of queries in your test set.
- “rank of first relevant document” is the position where the earliest relevant document appears.
- **Interpretation:**
 - MRR emphasizes *quick discoverability* of at least one relevant answer.
 - The higher the MRR, the closer relevant results appear to the top for most queries.

Evaluating the Retriever - MRR (Mean Reciprocal Rank)

- **Definition:** A metric that focuses on the rank of the *first* relevant document. For each query q :

$$\text{MRR} = \frac{1}{|Q|} \sum_{q \in Q} \frac{1}{\text{rank of first relevant document}}$$

- $|Q|$ is the total number of queries in your test set.
- “rank of first relevant document” is the position where the earliest relevant document appears.

Example

- Query A’s first relevant document is at rank 1 $\rightarrow$ reciprocal rank = $1/1 = 1.0$
- Query B’s first relevant document is at rank 4 $\rightarrow$ reciprocal rank = $1/4 = 0.25$
- $\text{MRR} = (1.0 + 0.25) / 2 = 0.625$

Evaluating the Retriever - MAP (Mean Average Precision)

- **Definition:** A deeper metric that looks at *every* relevant document's rank. For each query, you compute the *average precision (AP)* and then average across queries.

$$\text{MAP} = \frac{1}{|Q|} \sum_{q \in Q} \text{AP}(q)$$

- **AP:** The mean of the precision values computed at each rank where a relevant document is found.

$$\text{AP}(q) = \frac{1}{R_q} \sum_{k=1}^N [\text{Precision@}k \times \text{rel}(k)]$$

- R_q = total number of relevant documents for query q .
- $\text{rel}(k)$ is 1 if the document at rank k is relevant, else 0.
- Summation over all ranks k up to N (the number of retrieved documents or a cutoff).

Evaluating the Retriever - MAP (Mean Average Precision)

$$AP(q) = \frac{1}{R_q} \sum_{k=1}^N [\text{Precision}@k \times \text{rel}(k)]$$

Example (AP Calculation)

- Assume a query has 3 relevant documents in total.
- Suppose the ranked results are: [DocA (relevant), DocB (irrelevant), DocC (relevant), DocD (relevant), ...].
- The ranks where you hit relevant docs are ranks 1, 3, and 4.
 - $\text{Precision}@1 = 1/1 = 1.0$ (since DocA is relevant)
 - $\text{Precision}@3 = 2/3 \approx 0.667$ (2 relevant out of top 3)
 - $\text{Precision}@4 = 3/4 = 0.75$ (3 relevant out of top 4)
 - $AP = (1.0 + 0.667 + 0.75) / 3 \approx 0.806$
- Repeat for all queries and take the average for MAP.

Interpretation:

- High MAP suggests that a system *consistently* ranks relevant documents near the top across all queries.

Evaluating the Generator – ROUGE

ROUGE (Recall-Oriented Understudy for Gisting Evaluation)

Origins & Purpose:

- Developed primarily for text *summarization* tasks, focusing on how much of the important content in a reference summary is captured by the generated summary.

Key Variants:

- **ROUGE-N:** Measures overlap of *n-grams* between the generated text and reference. Commonly, ROUGE-1 (unigrams) and ROUGE-2 (bigrams) are used.
- **ROUGE-L:** Focuses on the *longest common subsequence (LCS)* between generated text and reference.

Focus on Recall:

- ROUGE often emphasizes “how much of the reference text’s key content is retained in the generated text.”
- Especially suitable when completeness of reference content is critical (e.g., summarization).

Typical Interpretation:

- **ROUGE-N:** Higher scores mean more overlap in n-gram usage between system output and reference.
- **ROUGE-L:** Higher scores mean the system output shares a longer subsequence pattern with the reference.

Evaluating the Generator – ROUGE Example

Reference Summary (R): "The cat sat on the mat"

Generated Summary (H): "The cat is on the mat"

ROUGE-1 (Unigram Overlap)

1. Identify Unigram Frequencies

Reference (R) Frequencies: Hypothesis (H) Frequencies:

"the": 2 "the": 2

"cat": 1 "cat": 1

"sat": 1 "is": 1

"on": 1 "on": 1

"mat": 1 "mat": 1

2. Calculate Overlap

For each unique token, overlap = min(freq in R, freq in H).

- "the": min(2, 2) = 2
- "cat": min(1, 1) = 1
- "sat": min(1, 0) = 0 (since "sat" does not appear in H)
- "on": min(1, 1) = 1
- "mat": min(1, 1) = 1
- "is": min(0, 1) = 0 (since "is" does not appear in R)

Total matching unigrams = 2 + 1 + 0 + 1 + 1 + 0 = 5.

3. ROUGE-1 Recall & Precision

Recall

= (Number of matching unigrams) / (Total unigrams in R)

= 5 / 6 $\approx$ 0.8333

Precision

= (Number of matching unigrams) / (Total unigrams in H)

= 5 / 6 $\approx$ 0.8333

4. ROUGE-1 F1

The F1 measure (or F-score) is the harmonic mean of Precision and Recall:

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

With both Precision and Recall = 0.8333,

$$F1 \approx 2 \times \frac{0.8333 \times 0.8333}{0.8333 + 0.8333} \approx 0.8333$$

Evaluating the Generator – ROUGE Example

Reference Summary (R): "The cat sat on the mat"

Generated Summary (H): "The cat is on the mat"

ROUGE-2 (Bigram Overlap)

1. List All Bigrams (consecutive pairs from the 6 tokens)

Reference Bigrams:

1. "the cat"
2. "cat sat"
3. "sat on"
4. "on the"
5. "the mat"

Hypothesis Bigrams:

1. "the cat"
2. "cat is"
3. "is on"
4. "on the"
5. "the mat"

2. Identify Common Bigrams

Compare each bigram in R to those in H.

- "the cat" → appears in H
- "cat sat" → **does not** appear in H
- "sat on" → **does not** appear in H
- "on the" → appears in H
- "the mat" → appears in H

Common bigrams = 3 ("the cat", "on the", "the mat").

3. ROUGE-2 Recall & Precision

Reference bigrams count = 5

Hypothesis bigrams count = 5

Overlap bigrams = 3

Recall = $3 / 5 = 0.6$

Precision = $3 / 5 = 0.6$

4. ROUGE-2 F1

With recall = precision = 0.6, F1 = 0.6.

Evaluating the Generator – ROUGE Example

Reference Summary (R): "The cat sat on the mat"

Generated Summary (H): "The cat is on the mat"

ROUGE-L (Longest Common Subsequence)

1. Longest Common Subsequence (LCS)

- ROUGE-L measures the length of the longest common subsequence (order matters, but tokens don't have to be consecutive).
- **Reference (R):** ["the", "cat", "sat", "on", "the", "mat"]
- **Hypothesis (H):** ["the", "cat", "is", "on", "the", "mat"]

One possible LCS is:

1. "the"
2. "cat"
3. "on"
4. "the"
5. "mat"

Here, the token "sat" in R and "is" in H are skipped, but the order of the remaining words is preserved. The length of this LCS is 5.

2. ROUGE-L Recall & Precision

Recall = LCS length / Reference length = $5 / 6 \approx 0.8333$

Precision = LCS length / Hypothesis length = $5 / 6 \approx 0.8333$

3. ROUGE-L F1

With recall = precision = 0.8333, F1 = 0.8333.

Evaluating the Generator – BLEU

BLEU (BiLingual Evaluation Understudy)

Origins & Purpose:

- First introduced for *machine translation*, comparing system translations to one or more reference translations.
- Now also applied to other text generation tasks (e.g., summarization, captioning) though it was initially geared toward translation.

Key Components:

- **N-gram Precision:** Checks how many n-grams in the system output appear in the reference.
- **Brevity Penalty:** Penalizes outputs that are *too short* compared to the reference, preventing a system from just returning short answers to avoid mistakes.

$$\text{BLEU} = \text{BP} \times \exp \left(\sum_{n=1}^N w_n \log p_n \right)$$

- p_n : precision of n-grams of length n .
- w_n : weight for n-grams (often uniform).
- **BP (Brevity Penalty):**

$$\text{BP} = \begin{cases} 1 & \text{if output_length} > \text{reference_length} \\ \exp(1 - \frac{\text{ref_length}}{\text{out_length}}) & \text{otherwise} \end{cases}$$

Evaluating the Generator – BLEU Example

Reference (R): “The cat sat on the mat”

Generated (H): “The cat is on the mat”

Unigram Precision (p1)

1. List unigrams

Reference: the (×2), cat (×1), sat (×1), on (×1), mat (×1)

Hypothesis: the (×2), cat (×1), is (×1), on (×1), mat (×1)

2. Count matches

- “the”: min(2 in R, 2 in H) = 2
- “cat”: min(1 in R, 1 in H) = 1
- “sat”: min(1 in R, 0 in H) = 0 (no “sat” in H)
- “on”: min(1 in R, 1 in H) = 1
- “mat”: min(1 in R, 1 in H) = 1
- “is”: min(0 in R, 1 in H) = 0 (no “is” in R)

Total matched unigrams = 2 + 1 + 1 + 1 = 5.

The hypothesis has 6 tokens total.

3. Precision@1 (p1)

$$p_1 = \frac{5}{6} \approx 0.8333$$

Bigram Precision (p2)

1. List bigrams

Reference bigrams (5 total):

- 1.(the, cat)
- 2.(cat, sat)
- 3.(sat, on)
- 4.(on, the)
- 5.(the, mat)

Hypothesis bigrams (5 total):

- 1.(the, cat)
- 2.(cat, is)
- 3.(is, on)
- 4.(on, the)
- 5.(the, mat)

2. Count bigram overlaps

- (the, cat) → match
- (cat, sat) vs. (cat, is) → no match
- (sat, on) vs. (is, on) → no match
- (on, the) → match
- (the, mat) → match

Total matched unigrams = 3.

The hypothesis has 5 bigrams total.

3. Bigram precision (p2)

$$p_2 = \frac{\text{matched bigrams}}{\text{total bigrams in hypothesis}} = \frac{3}{5} = 0.6$$

Evaluating the Generator – BLEU Example

BLEU (without smoothing) uses the geometric mean of the n-gram precisions.

Since we're only going up to 2-grams:

$$\text{mean_log_p} = \frac{1}{2} (\ln(p_1) + \ln(p_2)).$$

Let's plug in the values:

- $p_1 \approx 0.8333$
- $p_2 = 0.6$

$$\ln(0.8333) \approx -0.1823, \quad \ln(0.6) \approx -0.5108$$

$$\text{mean_log_p} = \frac{1}{2} ((-0.1823) + (-0.5108)) = \frac{-0.6931}{2} = -0.34655$$

$$\exp(\text{mean_log_p}) = \exp(-0.34655) \approx 0.7065$$

Evaluating the Generator – BLEU Example

The brevity penalty ensures that very short hypotheses (relative to references) aren't unfairly rewarded. It's given by:

$$\text{BP} = \begin{cases} 1 & \text{if hyp_length} \geq \text{ref_length} \\ \exp\left(1 - \frac{\text{ref_length}}{\text{hyp_length}}\right) & \text{otherwise} \end{cases}$$

Here:

- Reference length = 6
- Hypothesis length = 6

$$\text{ratio} = \frac{6}{6} = 1 \Rightarrow \text{BP} = 1.$$

Combine the geometric average of n-gram precisions with the brevity penalty:

$$\text{BLEU} = \text{BP} \times \exp\left(\frac{1}{2}(\ln(p_1) + \ln(p_2))\right).$$

$$\text{BLEU} = 1 \times 0.7065 = 0.7065$$

ROUGE vs BLEU

Recall vs. Precision Focus

- **ROUGE** highlights how much of the reference text's information the system output "recalls."
- **BLEU** highlights how many of the system output's phrases (n-grams) match the reference (i.e., *precision*).

Typical Use Cases

- **ROUGE**: Summarization evaluations, where capturing *all important points* from a reference is crucial.
- **BLEU**: Machine translation tasks, where matching the reference translation (especially at the phrase level) is key.

Limitations

- Both are *lexical* metrics, so synonyms or paraphrases can lower the score even if the meaning is correct.
- Neither metric directly measures *factual correctness* or deeper semantic equivalence.

Evaluating the Generator – BERTScore

Origins & Purpose:

- A more *semantically aware* metric using deep contextual embeddings from models like BERT.
- Compares each token's embedding in the hypothesis (generated text) to each token's embedding in the reference, measuring similarity in vector space.

Core Mechanism:

- **Token Embeddings:** Convert tokens in both reference and generated text into contextual embeddings (e.g., from BERT).
- **Similarity Computation:** Compute cosine similarity for each reference token vs. each generated token.
- **Precision / Recall & F1:**
 - BERTScore aligns each token in one text to the *most similar* token in the other text.
 - Summarizes these alignments into a precision, recall, or F1 measure at the *semantic* level.

Advantages:

- **Paraphrase Resilience:** Captures semantic similarity even if words differ lexically.
- **Better Correlation:** Tends to align more closely with human judgments for tasks involving paraphrasing or flexible wording.

Limitations:

- Still not perfect. Embedding models can be biased or incomplete for rare words.
- Performance can depend on the specific language model used (e.g., different BERT variants).

Evaluating the Generator – BERTScore

For each reference token r_i , find the hypothesis token h_j with the highest similarity. Sum those maxima over all reference tokens, normalize by the number of reference tokens $\rightarrow$ that's your *recall* side. Vice versa for the *precision* side (each hypothesis token finds its best-matching reference token). The final **BERTScore** is typically reported as an F1 measure combining those two.

$$\text{Precision} = \frac{1}{|H|} \sum_{h \in H} \max_{r \in R} \text{sim}(h, r), \quad \text{Recall} = \frac{1}{|R|} \sum_{r \in R} \max_{h \in H} \text{sim}(r, h)$$

$$\text{BERTScore}_{\text{F1}} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

RAG Evaluation – LLM as a judge

Why Use LLMs as Judges?

- **Scalability:** Human evaluations are often slow and expensive. An LLM-based judge can *rapidly* provide feedback on large volumes of text.
- **Consistency:** Large Language Models (LLMs) can apply the same *criteria* across different generations more consistently than multiple human raters with varying perspectives.
- **Versatility:** An LLM can be prompted to evaluate **various aspects**—fluency, factual accuracy, style compliance, etc.—depending on the task requirements.

LLM as a judge - Typical Setup

Prompt Construction:

- Provide the original query, the retrieved context (if any), and the generated text to the LLM judge.
- Include *evaluation instructions*, such as “Rate how factually correct this answer is on a scale from 1 to 5” or “Check if the answer is grounded in the provided text.”

LLM Evaluation Output:

- The LLM produces a textual explanation, numerical rating, or pass/fail verdict.
- It may also highlight *specific strengths and weaknesses* or *point out contradictions*.

LLM as a judge - Common Evaluation Dimensions

Factual Accuracy: Does the generated text align with the source material and common knowledge?

Coherence & Fluency: Is the text grammatically correct and logically structured?

Completeness: Does it address all key points from the prompt or reference?

Style & Tone: (If relevant) Are style guidelines (formal, casual, technical) followed?

Some RAG Evaluation Packages

- DeepEval: <https://github.com/confident-ai/deepeval>
- Ragas: <https://github.com/explodinggradients/ragas>
- TruLens: <https://github.com/truera/trulens>

LLM Application: Agents

What is an LLM Agent?

Definition:

- An LLM Agent is a system that uses a Large Language Model as its core reasoning engine, enhanced with additional capabilities such as planning, memory, and tool usage.

Key Components:

- **LLM Core:** The language model for understanding and generating text.
- **Planning:** Logic or strategy for sequencing actions.
- **Memory:** Storing context and historical information.
- **Tools:** External APIs or modules the agent can use.

GenAI Model

Which week is the recess week?

No context for internal policy.

What's the value of $(e^{2ix} + e^{-2ix})/4 + e^{2\ln(\sin(x))} + 1/2$?


Not good at complex calculations.

Who won the 2024 United States presidential election?

Knowledge cutoff.

Plan a trip for June holiday

Need to break down a complex task to a list of sub-tasks.



Which week is the recess week?

What's the value of $(e^{2ix} + e^{-2ix})/4 + e^{2\ln(\sin(x))} + 1/2$?

How's the 2024 United States presidential debates?

Plan a trip for June holiday

Planning

Introduction to Reasoning & Planning in LLM Agents

Planning allows agents to break down problems, decide which tools to use, and in which sequence. The interplay between LLM “reasoning” and explicit planning modules can significantly improve task performance.

- Planning is crucial for complex tasks
- Common patterns: linear plans, hierarchical decomposition, goal-based approaches
- Emergence of advanced prompting techniques (Chain-of-Thought, etc.)

Chain-of-Thought (CoT)

Chain-of-Thought (CoT) is best understood by comparing it to typical one-shot answers. If you ask a standard language model a question, it may provide a single direct response. With CoT, you prompt the model to explain each step it took mentally—like showing its “work” in a math problem. This can enhance interpretability and help identify mistakes early in the reasoning.

Definition:

- CoT is a prompting strategy that instructs a language model to generate intermediate reasoning steps before providing a final answer.

Purpose:


- Make the LLM’s reasoning process more transparent.
- Potentially improve accuracy and reduce “leaps” to incorrect conclusions.

Core Idea:

- Instead of jumping directly to the answer, the model “thinks out loud” in text, providing a step-by-step breakdown.

How Chain-of-Thought Works

Think of CoT prompting like guiding the model to emulate a teacher explaining every step to a student.



Magic prompt: Let's do this step by step...

Benefits & Challenges of CoT

Benefits:

- **Transparency:** can see how the model reaches its conclusions.
- **Improved Accuracy:** Breaking tasks down often leads to fewer mistakes for multi-step problems.
- **Error Debugging:** If a conclusion is incorrect, you can trace the faulty step.

Challenges:

- **Longer Output / Higher Token Usage:** CoT consumes more tokens, which can increase computation time and cost.
- **Still Susceptible to Hallucinations:** The model can produce a “logical-sounding” but flawed chain of reasoning.
- **Privacy Concerns (if used publicly):** CoT may reveal sensitive or proprietary reasoning steps.

Planning Approaches

Sequential (Linear) Planning

Sequential or linear planning is among the most straightforward planning methods. Think of it as a direct path from a starting point (e.g., a specific goal or question) to a final solution or outcome. In many everyday tasks—like following a recipe—this linear approach is both effective and efficient.

- **Definition:** A planning approach where tasks are executed one after another in a single, uninterrupted chain.
- **Characteristics:** Each step depends on the outcome of the previous step. There is typically a clear start, a definite end, and minimal branching or backtracking.
- **When to Use:** Straightforward tasks, or tasks with minimal branching factors.
- **Example:** Step-by-step guide to booking a flight or performing a simple calculation.

Hierarchical Planning

Hierarchical planning is particularly useful for complex or multi-faceted projects. Instead of treating an entire project as a single linear sequence, you break it down into a structured hierarchy. By creating this hierarchy, you can tackle each sub-goal with specialized strategies or tools, similar to a project manager assigning tasks to different team members.

Definition

- A planning paradigm that breaks a complex goal into multiple layers of sub-goals and tasks.

Hierarchy Levels

- **High-Level Goals:** The overarching mission or objective.
- **Sub-Goals:** Intermediate milestones.
- **Atomic Tasks:** The smallest actionable steps.

Why Hierarchical?

- Manage complexity by structuring large tasks into manageable units.
- Facilitates parallel work on sub-goals and potentially reusing sub-solutions.

Decomposing Goals into Sub-Goals & Tasks

A structured breakdown helps you see the entire scope and interdependencies. This is like the “divide and conquer” principle: each sub-goal becomes more approachable, and once the sub-goals are done, the entire system can be assembled.

Identify the High-Level Objective

- Example: “Build a web application” or “Write a research paper.”

Define Major Sub-Goals

- Example for a web app: Frontend, Backend, Database integration, Testing.
- Example for a research paper: Literature review, Data collection, Analysis, Drafting, Editing.

Drill Down into Atomic Tasks

- Each sub-goal can be decomposed further (e.g., for ‘Frontend’: design layout, code components, user testing).

Hierarchical Planning in Action: Research Paper Example

In this workflow, each sub-goal can be handled by a specialized “sub-agent” or module (or by a human user). For instance, an LLM-based agent could handle summarizing papers for the literature review or drafting a first pass. Once sub-goals are completed, the partial outputs integrate to form the entire paper.

High-Level Goal:

- “Write a 10-page research paper on renewable energy.”

Sub-Goals:

- Literature review
- Data gathering (statistics, case studies)
- Analysis & interpretation
- Drafting & editing
- Final proofread and formatting

Atomic Tasks:

- For “Literature review,” gather articles from academic databases; summarize key findings.
- For “Drafting,” structure the paper into an introduction, methodology, results, and conclusion.

ReACT (Reason+Act) Approach

In a traditional pipeline, an LLM might generate a chunk of reasoning and then present a final answer, without intermediate checks. **ReACT** ensures the model can pause mid-way, use tools, update its knowledge, and then continue reasoning. This iterative loop of “think → do → think → do” underlies more advanced use cases, like question answering with web searches or code execution.

Definition

- **ReACT** stands for “Reason + Act,” a strategy in which an LLM agent alternates between reasoning (thinking) and taking specific actions (e.g., calling tools) in a step-by-step manner.

Key Idea

- The model doesn’t just generate a final answer; it periodically “reflects” on new information or partial results, then decides the next action.

Motivation

- Combining reasoning and actions makes the agent more interactive, adaptable, and better at handling dynamic tasks.

The ReACT Workflow

ReACT can be viewed as a short feedback loop, repeated many times. Instead of planning everything in a single pass, the agent reevaluates each time it obtains new information. This approach is similar to how a human would solve a complex problem: think about what needs to be done, do it, see the result, and think again.

1. **Initial Prompt / Query:** The agent receives a user query or a task.
2. **Reasoning Phase:** The LLM “thinks out loud,” forming a partial plan or identifying information gaps.
3. **Action Phase:** The agent calls a tool or takes some external action (e.g., web search, database query).
4. **Observation / Feedback:** The agent receives results from the tool or environment.
5. **Refined Reasoning:** The LLM updates its plan based on the new data.
6. **Repeat or Final Answer:** This cycle continues until the agent converges on a final solution or response.

ReACT in Action: A Web Research Example

Notice how the agent doesn't just answer in one shot. It obtains relevant data mid-process, uses that data for further reasoning, and potentially repeats the cycle. This iterative method allows for more accurate and context-aware answers.

Task: "Find the cheapest flight from City A to City B next weekend and summarize the best option."

Step-by-Step ReACT

1. **Reason (Step 1):** The agent identifies it needs flight data.
 2. **Act (Tool Call):** The agent calls a flight search API.
 3. **Observe (Results):** Receives flight options and prices.
-
1. **Reason (Step 2):** The agent compares prices and connection times.
 2. **Act (Decision):** Chooses the best flight, checks if there's any discount code.
 3. **Final Answer:** Summarizes the recommendation for the user.

Building ReACT Agents

Implementation typically requires more than just a single LLM prompt. You'll likely have a loop in your code that updates the prompt with new evidence or tools' outputs, then re-invokes the LLM. Also consider how you handle partial information or incomplete data—ReACT truly shines when it can “try, observe, revise.”

Prompt Engineering

- Use instructions like “Think step by step. If you need information, call the appropriate tool.”

Tool Interfaces

- Provide the LLM with structured commands (e.g., “SEARCH_FLIGHTS(destination=XYZ)”) and parse the output.

State Management

- Maintain a record of the agent's last reasoning step, the results it observed, and what the next action might be.

Stopping Criteria

- Decide how many ReACT cycles are allowed or when a final answer is “good enough.”

Validation & Safety

- Check returned results for errors, and ensure the agent's actions remain within safe, intended boundaries.

Advantages & Challenges of ReACT

ReACT is great for real-time or dynamic tasks but demands more careful orchestration. Each iteration increases computational overhead, so you must weigh the benefits (adaptability, accuracy) against the cost (resource usage, potential complexity in design).

Advantages

- **Adaptive & Interactive:** The agent can change its plan if new information arises.
- **Reduced Hallucinations:** By verifying partial facts with tools, the agent may produce fewer fictitious details.
- **Transparent Thought Process:** Observers see how the agent reasons about each new piece of information.

Challenges

- **Complex Implementation:** Requires a system that can handle repeated queries to tools and keep track of state.
- **Higher Token & Time Cost:** Each reasoning step consumes tokens; multiple tool calls add latency.
- **Dependency on tools calling:** Heavily relying on tools calling. Will run into trouble when the agent making errors on how to call the tools.

Tools Use and Function Calling

Tools: Extending LLM Capabilities

Think of “tools” as the external services that your AI agent can call upon. An LLM alone might reason about a math problem, but by connecting it to a “calculator” tool, it can offload precise calculations. This extends the agent’s functionality and helps keep the LLM’s “internal reasoning” more accurate and efficient.

Definition

- In the context of LLM-based agents, “tools” are external applications or APIs (function calls) that the agent can invoke to extend its capabilities beyond text generation.

Why Tools Matter

- LLMs are great at generating and interpreting text, but they can’t inherently access real-time data or perform specialized computations.
- Tools allow the agent to retrieve up-to-date information, run code, process images, interact with databases, etc.

Examples of Common Tools

The exact tools your agent connects to depend on the domain. For instance, a coding assistant might need a live code runner, while a business analytics agent may rely on database queries and reporting APIs. Each tool can be wrapped in a function call or an endpoint the agent can invoke when necessary.

Web APIs

- Search engines, social media feeds, news sources for real-time information.

Database Queries

- Structured or unstructured data retrieval (SQL, NoSQL, vector databases).

Mathematical/Computational Tools

- Calculators for arithmetic, symbolic math libraries for advanced equations.

Code Execution Environments

- Python runners, JavaScript sandboxes for dynamic code generation and testing.

Domain-Specific APIs

- Finance (stock quotes, trading APIs), medical (drug interaction checkers), travel (flight/hotel booking).

Result Parsing & Validation

Even if a tool is reliable, mistakes can happen—like parsing errors or receiving partial data. The agent should be able to handle these gracefully. For example, if you ask an LLM agent to get a temperature reading from a weather API but the API returns an error, the agent should detect that the data is missing or corrupted, then attempt a retry or fallback logic.

Structured vs. Unstructured Outputs

- Tools may return data as JSON (structured) or raw text (unstructured).

Parsing Techniques

- For structured data: parse the format (JSON → Python dictionaries).
- For unstructured data: rely on prompts, regex, or further LLM-based summarization.

Validation & Sanity Checks

- Agents can verify if results are within expected ranges or match certain criteria.
- If results are invalid or contradictory, the agent can request a re-check or choose a different tool.

Memory

Introduction to Memory in LLM Agents

Although LLMs are inherently “stateless” in their raw form, adding memory mechanisms empowers them to handle multi-turn interactions, reference historical data, or even adapt based on user feedback. This makes them more flexible and context-aware over extended periods or sessions.

Definition

- “Memory” refers to how an LLM agent retains and recalls information over time—spanning from the current conversation (short-term) to potentially large repositories of past interactions or data (long-term).

Importance

- Enables continuity across multiple steps or conversations.
- Allows the agent to learn from previous actions, results, or user feedback.

Analogy to Human Memory

- Humans use working (short-term) memory to juggle immediate tasks.
- We store critical knowledge in long-term memory for later retrieval.

Short-Term Memory: Within the Context Window

Think of short-term memory as the immediate text the LLM “reads” every time you send a new prompt. Once you exceed the token limit, older content has to be truncated or summarized—if not done carefully, the agent might “forget” important details.

Definition

- Immediate context the LLM can see within its prompt, limited by the model’s context window

Usage

- Stores the ongoing conversation or the last few interactions.
- Keeps track of immediate tasks, sub-goals, or partial results.

Challenges

- Can overflow quickly if tasks or dialogues are long.
- Must carefully select what to keep in the prompt to avoid losing critical context.

Long-Term Memory: Beyond the Context Window

When you see an LLM recall something from a conversation that happened hours, days, or even months ago, it's often leveraging an external memory store. Upon receiving a new query, the agent can query this database using embeddings or keywords, retrieve the relevant info, and feed it back into the prompt—thus effectively “remembering” beyond the immediate context.

Definition

- An external store (e.g., database, vector embeddings) used to record past conversations, documents, or knowledge too large for the LLM's current prompt window.

Mechanism

- Often implemented via RAG.
- Use vector embedding search or other indexing to find relevant information from historical data.

Benefits

- Facilitates continuity over multiple sessions or large projects.
- Allows the agent to “remember” user preferences, previous solutions, or critical domain facts.

Agent Patterns

Multi-Agent Systems

Think of multi-agent systems as teams of LLM agents working together. Each one might have a distinct skill set (e.g., coding agent, math agent, summarization agent). When they communicate, share intermediate results, and coordinate strategies, you can tackle more complex or large-scale tasks than a single agent could handle efficiently.

Definition

- A **Multi-Agent System** (MAS) is composed of multiple interacting agents—each capable of perceiving, reasoning, and acting in an environment.

Why Multiple Agents?

- **Parallelism:** Divide a large task among specialized agents for speed or modularity.
- **Specialization:** Different agents can each have unique roles, domains, or toolsets.
- **Emergent Behavior:** Collaboration (and sometimes competition) can lead to solutions more powerful than a single agent approach.

Types of Multi-Agent Interactions

Multi-agent systems aren't strictly cooperative. In some scenarios (like trading bots), competition might drive better performance. In others (like generating an ensemble of solutions), collaboration is key. Understanding how agents align or diverge in goals influences system design significantly.

Collaborative Agents

- Share a common goal (e.g., working together to write a report).
- They exchange intermediate findings and solutions.

Competitive Agents

- Each agent pursues its own goal, which may conflict with others (e.g., game-playing agents).
- Can lead to strategic behavior or adversarial training.

Mixed-Mode Systems

- Some agents collaborate while others compete or remain neutral (e.g., negotiation or marketplace tasks).

Orchestrating Multiple Agents

The system design choice—centralized vs. distributed—often depends on the complexity of the task and how critical resilience or scalability is. For large-scale systems with many agents, distributed coordination can unlock better performance and redundancy.

Centralized Coordination

- A “manager” agent delegates sub-tasks to worker agents.
- Pros: Simpler to implement, easier to track.
- Cons: Single point of failure, potential bottleneck.

Distributed Coordination

- Agents share partial results or requests directly, using peer-to-peer messages.
- Pros: More robust, scalable.
- Cons: Requires more complex design to avoid deadlocks or conflicts.

Design Considerations for Multi-Agent Systems

Designing a multi-agent system involves orchestrating not just the communication flow but also the logic for conflict resolution. For instance, if you have multiple summarization agents producing slightly different outputs, you might rely on a “referee” agent to pick or merge the best one.

Agent Role Definition

- Clearly define each agent’s purpose and tools to avoid confusion.

Conflict Resolution

- If two agents produce conflicting answers, have a protocol for reconciling or escalating the conflict.

Consensus Mechanisms

- Voting systems or majority rules can help pick the best final solution from multiple proposals.

Scalability & Performance


- Too many agents can lead to excessive overhead.
- Use hierarchical “manager-employee” agent structures if needed.

Security & Privacy

- Agents might handle sensitive data. Ensure only the correct agent roles or individuals have access to that data.

Multi-Agent System Example

Objective: search for one potential company that might need green loans



Find Leads

Search for one potential large corporate client that might need green loans. Search for news where company talking about sustainability push or ESG initiatives that might require funding.

Focus the search for potential client in Asia. Skip all banks.

Generate a report that includes a comprehensive summary of the latest news, name of the company. Ensure that all source URLs are included in place in the report.

Business Research

Giving the potential lead, search for any mentions of company expansion or growth, mergers and acquisitions, infrastructure projects, market entry, privatization, de-listing, IPOs, government policies, and incentives that might suggest a need for a corporate loan.

Furthermore, research and identify existing banks that the company is currently working with.

Generate a Business Report that includes a comprehensive summary of the company business activities. Ensure that all source URLs are included inplace in the report.

ESG Research

Giving the potential lead, search for news articles discussing its sustainability initiatives, renewable energy projects, ESG (Environmental, Social, and Governance) strategies, and information on carbon emissions.

Generate an ESG Report that includes a comprehensive summary of the company's ESG related activities. Ensure that all source URLs are included in place in the report.

Generate Sales Report

Giving the lead, the ESG Report and the Business Report. Do additional research if needed to prepare a final sales report for the company as below.

Your final report must include the following sections:

Comprehensive Company Overview:

A brief description of the company and its business operations, including market capitalization, countries of operation, sector, major projects, managed assets, and decision-maker information (if available).

Business Activities:

Details of the company's business activities contained in the Business Report.

ESG Activities:

Details of the company's ESG activities and strategy contained in the ESG Report.

Finally, write the final report for the company to local file.

Task: Find Leads

Agent Name: Loan Researcher

Though: I need to search for recent news about companies in Asia that are discussing sustainability or ESG initiatives, as they might be potential clients for green loans. I will focus on the most recent news to ensure relevance.

Tool: Search news on the internet

Tool input: {"query": "Asia sustainability ESG initiatives 2024"}

Title: ESG's Future In Asia
Link: <https://gfmag.com/economics-policy-regulation/esg-in-asia/>
Snippet: The global sustainability trend has hit a snag the last two years. In the US, the movement to incorporate environmental, social and...

Title: Asia Pacific Real Estate Companies Aim to Achieve Net Zero
Link: https://www.csrwire.com/press_releases/797036-majority-asia-pacific-real-estate-companies-and-investment-funds-aim-achieve
Snippet: The role of Chief Sustainability Officer (CSO*) is rapidly gaining prominence among the Asia Pacific region as investors display increased...

Title: Asia Regulatory Update Newsletter December 2023
Link: <https://www.lw.com/en/insights/regulatory-updates-in-asia-esg-december-2023>
Snippet: In June 2023, the International Sustainability Standards Board (ISSB) released its inaugural sustainability disclosure standards,...

Title: Asian investors dial back ESG efforts but sustainability goals endure | ESG
Link: <https://www.asianinvestor.net/article/asian-investors-dial-back-esg-efforts-but-sustainability-goals-endure/492805>
Snippet: Survey shows clear slowdown amid challenging macro environment. Meanwhile, earlier initiatives are taking time to settle.

Task: Find Leads

Agent Name: Loan Researcher

Thought: The search results are **not very useful** in finding a **specific large Asian company**. I need to **specifically identify a large corporate client** in Asia that is actively pursuing sustainability or ESG initiatives in 2024.

Tool: Search news on the internet

Tool input: {"query": " large corporate clients in Asia sustainability ESG initiatives 2024"}

Title: Asia advances sustainability-related finance taxonomies, disclosure rules & public-private partnerships in ESG

Link: <https://www.thomsonreuters.com/en-us/posts/esg/asia-sustainability-rules/>

Snippet: Major developments in ESG policies in Asia indicate that ongoing progress in sustainability will continue throughout the region in 2024.

Title: ESG in APAC: 3 Trends to Watch in 2024

Link: <https://www.whitecase.com/insight-alert/esg-apac-3-trends-watch-2024>

Snippet: Faced with increased scrutiny by consumers, regulators, shareholders and boards, Environmental, Social and Governance (ESG) considerations...

Title: 35% of Asia/Pacific* Firms to Elevate ESG Reporting with AI by 2024

Link: <https://www.idc.com/getdoc.jsp?containerId=prAP51825324>

Snippet: IDC Reveals Sustainability/ESG Predictions for 2024 and Beyond. SINGAPORE, 30 January 2024 – In a recent report, IDC FutureScape: Worldwide...

Title: Companies walk ESG tightrope, under fire from all sides over disclosures

Link: <https://www.reuters.com/sustainability/boards-policy-regulation/companies-walk-esg-tightrope-under-fire-all-sides-over-disclosures-2024-02-29/>

Snippet: Since joining **Asahi Group in 2020 as sustainability chief, Preeti Srivastav has helped the Japanese brewer make several environmental** and...

Task: Find Leads

Agent Name: Loan Researcher

Thought: From the gathered results, I see mentions of several companies and initiatives. After analyzing the search results, I think I have found a promising lead.

Comprehensive Report:

Asahi Group, a prominent player in the beverage industry, has been making significant strides in sustainability and environmental stewardship. The company has been focusing on reducing its carbon footprint, improving energy efficiency, and adopting sustainable practices across its operations. Under the leadership of their sustainability chief, Preeti Srivastav, Asahi Group has implemented several initiatives aimed at enhancing their environmental performance.

Given their robust push towards sustainability, Asahi Group is a prime candidate for green loans which could further support their ESG initiatives and help them achieve their environmental goals.

Source URLs:

- [Reuters Article on Asahi Group's Sustainability Efforts](#)

This detailed analysis should impress the customer by showcasing the identification of a viable corporate client that might require green loans due to their ongoing sustainability initiatives.

Task: ESG Research

Agent Name: ESG Analyst

Thought: To create a comprehensive ESG report for Asahi Group, I need to gather detailed information about their recent sustainability initiatives, renewable energy projects, ESG strategies, and carbon emissions. First, I will search for recent news articles to get the most up-to-date information.

Tool: Search news on the internet

Tool Input: {"query": "Asahi Group sustainability initiatives 2024"}

Title: Asahi Group Holdings FY2024 Q1 Financial Results | Newsroom
Link: <https://www.asahigroup-holdings.com/en/newsroom/detail/20240514-0201.html>
Snippet: Asahi Group Holdings, Ltd. (Tokyo, Japan - May 14, 2024) – Asahi Group Holdings, Ltd. announced the financial results of the first quarter...

Title: Asahi Kasei starts operation of multi-module hydrogen pilot plant in Kawasaki | 2024 | News | Asahi Kasei
Link: <https://www.asahi-kasei.com/news/2024/e240514.html>
Snippet: The Japanese technology company Asahi Kasei is further accelerating its hydrogen business activities.

Title: Companies walk ESG tightrope, under fire from all sides over disclosures
Link: <https://www.reuters.com/sustainability/boards-policy-regulation/companies-walk-esg-tightrope-under-fire-all-sides-over-disclosures-2024-02-29/>
Snippet: Since joining Asahi Group in 2020 as sustainability chief, Preeti Srivastav has helped the Japanese brewer make several environmental and...

Title: Asahi Kasei to Highlight Sustainability, 3D Printing Filaments, Purging Compounds, and More at NPE 2024
Link: <https://www.businesswire.com/news/home/20240410419606/en/Asahi-Kasei-to-Highlight-Sustainability-3D-Printing-Filaments-Purging-Compounds-and-More-at-NPE-2024>
Snippet: APNA, Asaclean, and Elastomer will highlight new technologies at NPE 2024 from May 6 through 10 in Orlando, Florida.

Task: ESG Research

Agent Name: ESG Analyst

Thought: I now know the final answer

Comprehensive Summary

Asahi Group, a leading company in the beverage industry, has made significant progress in various sustainability initiatives. They have emphasized reducing their carbon footprint, improving energy efficiency, and adopting sustainable practices under the guidance of their sustainability chief, Preeti Srivastav.

One of the notable projects is the Asahi Forest, which focuses on preserving forest ecosystems while educating future generations about nature's importance. Additionally, Asahi Kasei, a subsidiary, is advancing in renewable energy with the operation of a multi-module hydrogen pilot plant in Kawasaki, marking a significant step towards green hydrogen production.

Further demonstrating their commitment, Tongsoh Petrochemical, another subsidiary, received a Gold Medal in the EcoVadis Sustainability Assessment, highlighting their achievements in sustainable practices.

These initiatives and recognitions underscore Asahi Group's dedication to environmental stewardship and position them as a prime candidate for green loans, which could further support their ESG goals.

Source URLs:

- [Asahi Group Holdings FY2024 Q1 Financial Results](#)
- [Asahi Kasei starts operation of multi-module hydrogen pilot plant in Kawasaki](#)
- [Companies walk ESG tightrope, under fire from all sides over disclosures](#)
- [Asahi Kasei to Highlight Sustainability, 3D Printing Filaments, Purging Compounds, and More at NPE 2024](#)
- [The Asahi Forest: A Natural Classroom, Bringing Nature's Gifts to Future Generations](#)
- [Asahi Kasei's subsidiary Tongsoh Petrochemical receives Gold Medal in EcoVadis Sustainability Assessment](#)
- [Asahi Kasei Unveils Hydrogen Pilot Plant in Kawasaki, Pioneering Green Hydrogen Production](#)

Task: Business Research

Agent Name: Loan Researcher

Though: Given the task at hand, I need to gather information on Asahi Group regarding recent corporate activities such as expansion, mergers, acquisitions, infrastructure projects, market entry, privatization, de-listing, IPOs, government policies, and incentives that might suggest a need for a corporate loan. Additionally, I need to identify the existing banks that Asahi Group is currently working with. First, I will search for recent news about Asahi Group to gather relevant information on their corporate activities.

Tool: Search news on the internet

Tool Input: {"query": "Asahi Group recent corporate news 2024"}

Title: Asahi Group Holdings First Quarter 2024 Earnings: Beats Expectations

Link: <https://simplywall.st/stocks/jp/food-beverage-tobacco/tse-2502/asahi-group-holdings-shares/news/39d2f7f6d04ce51b>

Snippet: Asahi Group Holdings (TSE:2502) First Quarter 2024 Results Key Financial Results Revenue: JP¥616.6b (up 11% from 1Q...

Title: Asahi CEO Seeking Acquisition Targets to Expand Super Dry in US

Link: <https://www.bloomberg.com/news/articles/2024-02-22/asahi-ceo-seeking-acquisition-targets-to-expand-super-dry-in-us>

Snippet: Asahi Group Holdings Ltd. is ready to find more acquisition targets as the Super Dry beer brewer seeks to push deeper into the US,...

Title: Asahi Kasei to Construct a Lithium-ion Battery Separator Plant in Canada | 2024 | News | Asahi Kasei

Link: <https://www.asahi-kasei.com/news/2024/e240425.html>

Snippet: Asahi Kasei announced today that it will construct an integrated plant in Ontario, Canada for the base film manufacturing and coating of...

Title: Companies walk ESG tightrope, under fire from all sides over disclosures

Link: <https://www.reuters.com/sustainability/boards-policy-regulation/companies-walk-esg-tightrope-under-fire-all-sides-over-disclosures-2024-02-29/>

Snippet: Since joining Asahi Group in 2020 as sustainability chief, Preeti Srivastav has helped the Japanese brewer make several environmental and...

Task: Business Research

Agent Name: Loan Researcher

Thought: I have identified several key pieces of recent news related to Asahi Group's corporate activities, including acquisitions, expansions, infrastructure projects, and restructuring. Now, I need to gather information on which banks Asahi Group is currently working with.

Action: Search the internet

Action Input: {"query": "Asahi Group current banking partners 2024"}

Title: Asahi Group Holdings FY2024 Q1 Financial Results | Newsroom
Link: <https://www.asahigroup-holdings.com/en/newsroom/detail/20240514-0201.html>
Snippet: (Tokyo, Japan - May 14, 2024) – Asahi Group Holdings, Ltd. announced the financial results of the first quarter in 2024, for the period of ...

Title: Investors | ASAHI GROUP HOLDINGS
Link: <https://www.asahigroup-holdings.com/en/ir/>
Snippet: Investor relations information for shareholders and investors of the Asahi Group, including recent ... FY2024 Q1 Financial Results Presentation. FY2024 ...

Title: IR News | Investors | ASAHI GROUP HOLDINGS
Link: <https://www.asahigroup-holdings.com/en/ir/news/>
Snippet: Latest ; May 14, 2024 · FY2024 Q1 Financial Results ; May 14, 2024 · Monthly Sales Results (Apr.2024) ; May 08, 2024 · On-demand streaming of the 100th Annual General ...

Title: Asahi Group Announces New Group Corporate Logomark and ...
Link: <https://www.asahigroup-holdings.com/en/newsroom/detail/20240305-0201.html>
Snippet: proudly announced today the introduction of a new group corporate logomark and corporate statement, set to take effect from April 1, 2024.

Task: Business Research

Agent Name: Loan Researcher

Thought: I now know the final answer

Business Report on Asahi Group

Comprehensive Summary of Business Activities

Expansion and Growth:

1. Acquisition Targets in the US:

- **Title:** Asahi CEO Seeking Acquisition Targets to Expand Super Dry in US
- **Link:** [Bloomberg](#)
- **Snippet:** Asahi Group Holdings Ltd. is ready to find more acquisition targets as the Super Dry beer brewer seeks to push deeper into the US.

2. Sales Surge:

- **Title:** Asahi Super Dry sales see 35% volume surge
- **Link:** [The Drinks Business](#)
- **Snippet:** Total sales of Asahi Super Dry beer soared by 35% in 2023 as the brand's global expansion plans pay off.

3. US Market Push:

- **Title:** With U.S. push, Asahi aims to sate evolving tastes
- **Link:** [Japan Times](#)
- **Snippet:** Asahi Group Holdings is pushing into the crowded U.S. beverage market, placing its “premium” offerings and low-alcohol and nonalcoholic beverages.

Task: Business Research

Agent Name: Loan Researcher

Thought: I now know the final answer

Mergers and Acquisitions:

1. New Acquisition:

- **Title:** Well Wishes to Isaac! - Octopi Brewing Has Been Acquired By Asahi
- **Link:** [The Full Pint](#)
- **Snippet:** Asahi Group has acquired Octopi Brewing, a strategic move to strengthen its brewing capabilities and market presence.

Infrastructure Projects:

1. New Plant in Canada:

- **Title:** Asahi Kasei to Construct a Lithium-ion Battery Separator Plant in Canada
- **Link:** [Asahi Kasei](#)
- **Snippet:** Asahi Kasei announced today that it will construct an integrated plant in Ontario, Canada for the base film manufacturing and coating of lithium-ion battery separators.

Corporate Restructuring:

1. Group Restructuring:

- **Title:** 'Maximising corporate value': Asahi hopes massive group restructure will enhance advantages across four key markets
- **Link:** [Food Navigator Asia](#)
- **Snippet:** Japanese beer giant Asahi is hopeful that its recently-completed large-scale group restructuring strategy to establish individual regional entities will improve its market positioning.

Task: Business Research

Agent Name: Loan Researcher

Thought: I now know the final answer

Financial Performance:

1. Q1 2024 Financial Results:

• Title: Asahi Group Holdings First Quarter 2024 Earnings: Beats Expectations

• Link: [Simply Wall St](#)

• Snippet: Asahi Group Holdings (TSE:2502) First Quarter 2024 Results Key Financial Results Revenue: JP¥616.6b (up 11% from 1Q...

Existing Banking Relationships

1. DBS Bank:

• Title: Asahi Group Holdings Ltd - DBS Bank

• Link: [DBS Bank](#)

• Snippet: Asahi Group Holdings Ltd. The latest investment analysis on Asahi Group Holdings Ltd. Group Research - Equities21 Mar 2024.

2. Nishimura & Asahi:

• Title: Banking | Experience | Nishimura & Asahi

• Link: [Nishimura & Asahi](#)

• Snippet: Nishimura & Asahi Receives Top Rankings in Chambers Global 2024 · 2024.2. Outstanding Results for Nishimura & Asahi at the Legal 500 Asia Pacific 2024 · 2023.12.

Source URLs:

• [Reuters Article on Asahi Group's Sustainability Efforts](#)

• [Bloomberg Article on Asahi CEO Seeking Acquisition Targets](#)

• [Asahi Kasei News on New Plant in Canada](#)

• [Food Navigator Asia Article on Asahi Restructure](#)

• [The Full Pint Article on Octopi Acquisition](#)

• [Japan Times Article on US Market Push](#)

• [The Drinks Business Article on Sales Surge](#)

Task: Generate Sales Report

Agent Name: Treasury Product Sales

Thought: To prepare the final sales report for Asahi Group, I will need to gather comprehensive information about the company, its business activities, and its ESG activities.

Here's my plan:

1. Perform a search on the internet to get a comprehensive overview of Asahi Group, including its market capitalization, countries of operation, sector, major projects, managed assets, and decision-maker information.
2. Search for recent news articles related to Asahi Group to ensure the information is up-to-date.
3. Summarize the business report and ESG report provided.
4. Compile all the information into a final report and save it to a local file.

Let's begin!

Tool: Search the internet

Tool Input: {"query": "Asahi Group company overview 2024"}

Title: Asahi Group Holdings FY2024 Q1 Financial Results | Newsroom
Link: <https://www.asahigroup-holdings.com/en/newsroom/detail/20240514-0201.html>
Snippet: (Tokyo, Japan - May 14, 2024) – Asahi Group Holdings, Ltd. announced the financial results of the first quarter in 2024, for the period of ...

Title: Asahi Group Announces New Group Corporate Logomark and ...
Link: <https://www.asahigroup-holdings.com/en/newsroom/detail/20240305-0201.html>
Snippet: proudly announced today the introduction of a new group corporate logomark and corporate statement, set to take effect from April 1, 2024. This ...

Title: [PDF] Asahi Group Holdings, Ltd. Q1 2024 Financial Results Briefing
Link: https://www.asahigroup-holdings.com/pdf/en/ir/event/presentation/2024_1q_presentation_summary_1.pdf
Snippet: Asahi Group Holdings, Ltd. Q1 2024 Financial Results Briefing: Summary of Conference Call (1). Date: Tuesday, May 14, 2024, 18:00-19:00.

Title: Asahi Group Holdings Company Profile 2024 - PitchBook
Link: <https://pitchbook.com/profiles/company/59611-06>
Snippet: Asahi Group Holdings General Information. Description. Asahi is a leading brewer in Japan with an estimated 36.5% market share by volume in Japan's beer and ...

Final Sales Report for Asahi Group

Comprehensive Company Overview

Asahi Group Holdings, Ltd. is a leading brewer in Japan, with a significant market share in the beer and beverage market. The company operates globally, with key markets in Japan, Europe, Australia, and Southeast Asia. Recent major projects include significant acquisitions and restructuring efforts aimed at global expansion and sustainability. The company is led by a proactive CEO and President who is focused on further globalization and sustainability.

Key Information

- **Name:** Asahi Group Holdings, Ltd.
- **Website:** [Asahi Group Holdings](#)
- **Sector:** Food and Beverage
- **Market Capitalization:** As of Q1 2024, Asahi Group is a leading brewer in Japan with an estimated 36.5% market share by volume in Japan's beer and beverage market.
- **Countries of Operation:** Japan, Europe, Australia, and Southeast Asia.
- **Major Projects:** Acquisitions, restructuring efforts, and sustainability initiatives.
- **Managed Assets:** Diverse portfolio of alcoholic beverages, soft drinks, and food products.
- **Decision-Makers:** Asahi Group CEO and President are actively involved in driving the company's globalization and sustainability strategies.

Business Activities

Asahi Group's business operations cover alcoholic beverages, soft drinks, and food products. The company has been actively involved in corporate restructuring and expansion, including the acquisition of Never Never, a gin maker.

Detailed Business Activities:

1. **Alcoholic Beverages:** Manufacturing and sale of beer, spirits, and other alcoholic drinks. This segment contributes significantly to the company's revenue.
2. **Soft Drinks:** Production and distribution of non-alcoholic beverages, including carbonated drinks, juices, and bottled water.
3. **Food Products:** The company is involved in the production and sale of food items, catering to various consumer needs.
4. **Corporate Changes/Expansions:** In 2024, Asahi Group announced significant restructuring to enhance its competitive advantages across key markets. The company has also been actively acquiring other brands, such as the recent acquisition of Never Never, a gin maker.

Sources:

- [Asahi Group Holdings FY2024 Q1 Financial Results](#)
- [Asahi Group Announces New Group Corporate Logomark](#)
- [Asahi Group Holdings Company Profile](#)

ESG Activities

Asahi Group is committed to sustainability, with significant investments in renewable energy projects and initiatives to reduce carbon emissions. The company has set ambitious targets to achieve carbon neutrality and is actively working on sustainable sourcing and waste reduction.

Detailed ESG Activities:

1. **Renewable Energy Projects:** Asahi Group has invested significantly in renewable energy sources to power their operations, aiming to reduce reliance on fossil fuels.
2. **Sustainability Initiatives:** The company has implemented various programs to ensure sustainable sourcing of ingredients, reduce waste, and promote recycling.
3. **Carbon Emissions:** Asahi Group is committed to reducing its carbon emissions and has set ambitious targets to achieve carbon neutrality in the coming years.

Sources:

- Comprehensive ESG Report for Asahi Group (Internal Document)
- Recent news articles supporting Asahi Group's commitment to sustainability ([Asahi Group Holdings First Quarter 2024 Earnings](#))

Reasoning Model as Planner

You're the most experienced loan products sales in a bank. You are tasked to search for one potential company that might need green loans. You have Google News search tools. Generate a detailed action plan for this task.

Action Plan: Identifying Potential Companies for Green Loans Using Google News

1. Define Target Industries & Keywords

- **Industries:** Renewable energy, clean tech, sustainable agriculture, green construction, manufacturing (eco-friendly practices), waste management, EVs, and ESG-focused sectors.
- **Keywords:**
 - Project-related: "renewable energy project," "sustainability initiative," "carbon neutrality plan," "green building certification," "circular economy."
 - Financial needs: "seeking green financing," "sustainable investment," "funding for eco-project," "green bond issuance."
 - Regulatory: "new environmental regulation," "carbon tax compliance," "emissions reduction target."
 - Partnerships: "partnered with [Green Tech Firm]," "collaboration with UNEP/WWF."

Reasoning Model as Planner

You're the most experienced loan products sales in a bank. You are tasked to search for one potential company that might need green loans. You have Google News search tools. Generate a detailed action plan for this task.

Action Plan: Identifying Potential Companies for Green Loans Using Google News

2. Set Up Advanced Google News Search Tools

- **Filters:** Use location (e.g., "region:Asia "), time (past 3–6 months), and language.
- **Alerts:** Create Google Alerts for combinations like:
 - "solar farm" + "financing"
 - "net-zero target" + "manufacturing sector"
 - "LEED certification" + "new facility"
- **Operators:** Exclude irrelevant terms (e.g., -"grant" -"subsidy") and use quotes for exact phrases.
- Combine keywords using Boolean operators. For example:

arduino

 Copy

"renewable energy" AND ("investment" OR "project") AND "green loan"

Reasoning Model as Planner

You're the most experienced loan products sales in a bank. You are tasked to search for one potential company that might need green loans. You have Google News search tools. Generate a detailed action plan for this task.

Action Plan: Identifying Potential Companies for Green Loans Using Google News

3. Analyze News Articles for Triggers

- **Project Announcements:** Look for phrases like *"breaking ground on a new solar plant"* or *"launching a waste-to-energy program."*
- **Regulatory Compliance:** Companies facing new environmental laws (e.g., *"comply with EU Taxonomy regulations"*).
- **Awards/Certifications:** Recognition like *"CDP Climate Leader"* or *"RE100 member."*
- **Leadership Changes:** New sustainability officers or green-focused board appointments.
- **Financial Signals:** Mentions of capital-raising, partnerships with development banks, or delays in project funding.

Reasoning Model as Planner

You're the most experienced loan products sales in a bank. You are tasked to search for one potential company that might need green loans. You have Google News search tools. Generate a detailed action plan for this task.

Action Plan: Identifying Potential Companies for Green Loans Using Google News

4. Shortlist & Prioritize Companies

- **Alignment:** Prioritize companies with clear, measurable green projects (e.g., *"20MW wind farm by 2025"*).
- **Financial Health:** Cross-reference with platforms like Bloomberg or Crunchbase for revenue trends and debt ratios.
- **Competitor Activity:** If a competitor bank is mentioned, propose refinancing with better terms.

Reasoning Model as Planner

You're the most experienced loan products sales in a bank. You are tasked to search for one potential company that might need green loans. You have Google News search tools. Generate a detailed action plan for this task.

Action Plan: Identifying Potential Companies for Green Loans Using Google News

5. Deep-Dive Research

- **Sustainability Reports:** Check if the company publishes ESG reports (indicating structured commitments).
- **Supply Chain:** Suppliers to Fortune 500 companies with net-zero pledges may need green loans for upgrades.
- **Social Media:** Follow company LinkedIn/Twitter for real-time updates on unmet financing needs.

Reasoning Model as Planner

You're the most experienced loan products sales in a bank. You are tasked to search for one potential company that might need green loans. You have Google News search tools. Generate a detailed action plan for this task.

Action Plan: Identifying Potential Companies for Green Loans Using Google News

6. Outreach Strategy

- **Identify Decision-Makers:** Use news mentions (e.g., *"CFO Jane Doe announced..."*) or LinkedIn searches for sustainability roles.
- **Tailored Pitch:** Reference specific news (e.g., *"Congratulations on your new biogas initiative—we can structure a loan aligned with your 2026 goals"*).
- **Leverage Timing:** Contact within 2 weeks of a project announcement when budgeting is likely underway.

Reasoning Model as Planner

You're the most experienced loan products sales in a bank. You are tasked to search for one potential company that might need green loans. You have Google News search tools. Generate a detailed action plan for this task.

Action Plan: Identifying Potential Companies for Green Loans Using Google News

7. Validate & Mitigate Risks

- **Greenwashing Check:** Verify certifications (e.g., B Corp, ISO 14001) or partnerships with credible NGOs.
- **Regulatory Risk:** Ensure projects comply with local/international green standards (e.g., EU Taxonomy).

Reasoning Model as Planner

You're the most experienced loan products sales in a bank. You are tasked to search for one potential company that might need green loans. You have Google News search tools. Generate a detailed action plan for this task.

Action Plan: Identifying Potential Companies for Green Loans Using Google News

Tools Beyond Google News:

- **LinkedIn Sales Navigator:** Track employee updates.
- **ESG Databases:** Sustainalytics, MSCI.
- **Regulatory Trackers:** IISD SDG Knowledge Hub.

Autonomous Agents vs Agentic Workflow

Autonomous Agents

- To handle complexity and uncertainty by “thinking” for themselves rather than strictly following pre-written instructions.
- **Independence:** Operate without constant human intervention.
- **Reactivity:** Sense changes in the environment and adjust behavior accordingly.
- **Adaptability:** Learn from experiences and modify strategies in real time.
- **Flexibility:** Can handle dynamic, uncertain, or complex environments.
- **Innovation:** Can develop unexpected strategies through learning and adaptation.
- **Unpredictability:** Behavior may be less transparent or harder to predict.

Agentic Workflow

- **Predefined Workflow** agents execute tasks based on a rigid, pre-scripted sequence of steps. Their behavior is determined by a fixed process that has been designed and tested ahead of time.
- **Determinism:** Follow a strict, predictable path with little to no deviation.
- **Modularity:** Processes are broken down into clearly defined, discrete steps.
- **Predictability:** Outcomes and actions are easily anticipated.
- **Reliability:** Predictable behavior makes it easier to test and debug.
- **Efficiency:** Works very well in environments where tasks and outcomes are well understood.
- **Inflexibility:** Struggles with unexpected changes or novel situations.